

# Introduction to Multithreaded Programming in Java

Affan Rauf

National University of Computer and Emerging Sciences

December 18, 2024

# What is a Thread?

- ▶ A thread is the smallest unit of execution in a program.
- ▶ Threads allow a program to perform multiple tasks simultaneously.
- ▶ Threads within a process share memory, open files, and other resources.

# Overview of Memory Segments in a Program

- ▶ When a C program is loaded into memory, it is divided into four main segments:
  - ▶ **Code Segment (Text Segment):** Contains the executable code.
  - ▶ **Data Segment:** Holds global and static variables initialized at compile time.
  - ▶ **Heap Segment:** Used for dynamically allocated memory during runtime.
  - ▶ **Stack Segment:** Stores local variables, function calls, and return addresses.

# Memory Sharing Among Threads

- ▶ Not all memory segments are shared between threads. Here's how they differ:

Memory Segment	Shared by Threads?
Code Segment	Yes
Data Segment	Yes
Heap Segment	Yes
Stack Segment	No

# Advantages Of This Memory Design?

- ▶ **Code Sharing:**

- ▶ Reduces memory usage by allowing threads to share instructions.

- ▶ **Data/Heap Sharing:**

- ▶ Enables threads to collaborate on shared data structures.

- ▶ **Private Stacks:**

- ▶ Ensures independent function calls and local variable usage without conflicts.

# Extending Thread Class

- One way to create a thread is by extending the Thread class.

```
1  class MyThread extends Thread {
2      @Override
3      public void run() {
4          for (int i = 0; i < 5; i++) {
5              System.out.println("Thread running: " + i);
6              try {
7                  Thread.sleep(500); // Pause for 500 milliseconds
8              } catch (InterruptedException e) {
9                  System.out.println("Thread interrupted: " + e.getMessage());
10             }
11         }
12     }
13 }
14 public class ThreadExample {
15     public static void main(String[] args) {
16         MyThread thread = new MyThread();
17         thread.start();
18
19         for (int i = 0; i < 5; i++) {
20             System.out.println("Main thread: " + i);
21             try {
22                 Thread.sleep(500);
23             } catch (InterruptedException e) {
24                 System.out.println("Main thread interrupted: " + e.getMessage());
25             }
26         }
27     }
28 }
```

# Implementing Runnable Interface

- ▶ Another way to create a thread is by implementing the Runnable interface.

```
1  class MyRunnable implements Runnable {
2      @Override
3      public void run() {
4          for (int i = 0; i < 5; i++) {
5              System.out.println("Runnable running: " + i);
6              try {
7                  Thread.sleep(500); // Pause for 500 milliseconds
8              } catch (InterruptedException e) {
9                  System.out.println("Runnable interrupted: " + e.getMessage());
10             }
11         }
12     }
13 }
14 public class RunnableExample {
15     public static void main(String[] args) {
16         MyRunnable myRunnable = new MyRunnable();
17         Thread thread = new Thread(myRunnable);
18         thread.start();
19         for (int i = 0; i < 5; i++) {
20             System.out.println("Main thread: " + i);
21             try {
22                 Thread.sleep(500);
23             } catch (InterruptedException e) {
24                 System.out.println("Main thread interrupted: " + e.getMessage());
25                 ;
26             }
27         }
28     }
```

# What is Thread Interleaving?

- ▶ Threads in a multithreaded program can execute in an arbitrary order.
- ▶ This behavior is known as **thread interleaving**.
- ▶ Interleavings are determined by the JVM scheduler and the underlying operating system.
- ▶ Without proper synchronization, interleaving can lead to **race conditions**.



# Example: Incrementing a Shared Counter

## Scenario:

- ▶ Two threads increment a shared counter 1,000 times each.
- ▶ Shared variable: `int count = 0;`

## Code:

```
1  class Counter {  
2      int count = 0;  
3      public void increment() {  
4          count++;  
5      }  
6  }
```

**Expected Result:** `count = 2000` **Actual Result:** Can be less than 2000 due to interleaving!

# Breaking Down count++

## Steps in count++:

1. Read the value of count.
2. Increment the value.
3. Write the updated value back to count.

## Java Bytecode for count++:

```
1  load count    // Read count from memory
2  increment     // Increment value
3  store count   // Write back to memory
```

**Problem:** These steps are **not atomic**, meaning they can be interleaved.

# Possible Thread Interleavings

## Example Interleaving: Two Threads (T1 and T2)

Incrementing count

### Expected Execution:

1. T1: Read count = 0.
2. T1: Increment to 1.
3. T1: Write 1.
4. T2: Read count = 1.
5. T2: Increment to 2.
6. T2: Write 2.

### Actual Interleaving:

1. T1: Read count = 0.
2. T2: Read count = 0.
3. T1: Increment to 1.
4. T2: Increment to 1.
5. T1: Write 1.
6. T2: Write 1.

**Result:** Final count is 1, not 2!

# Why Does This Happen?

- ▶ The `count++` operation is not atomic (composed of multiple steps).
- ▶ Threads interleave arbitrarily between these steps.
- ▶ Updates to `count` are overwritten due to simultaneous reads and writes.
- ▶ This is called a **race condition**.

# Fixing the Problem: Synchronization

## Solution: Make `count++` **Atomic**

- ▶ Use the `synchronized` keyword to ensure only one thread executes `increment()` at a time.

## Fixed Code:

```
1  class Counter {  
2      int count = 0;  
3      public synchronized void increment() {  
4          count++;  
5      }  
6  }
```

**Result:** Correct final count = 2000.

# Need for Synchronization

- ▶ Multiple threads accessing shared resources can cause issues (e.g., race conditions).
- ▶ Java provides `synchronized` keyword for thread safety.

# What is synchronized?

- ▶ The `synchronized` keyword is used to ensure thread safety in Java.
- ▶ It prevents multiple threads from executing a critical section of code simultaneously.
- ▶ Achieves **mutual exclusion** by using a monitor (or intrinsic lock).

## Key Benefits

- ▶ Avoids **race conditions**.
- ▶ Protects **shared mutable resources**.
- ▶ Ensures **predictable behavior** in multithreaded programs.

# How It Works

1. A thread entering a synchronized block or method acquires the **monitor lock**.
2. Other threads attempting to enter the same synchronized block/method are **blocked**.
3. The thread releases the lock upon exiting the synchronized block or method.

## Monitor Lock

Each object in Java has a monitor lock that can be acquired or released.



# Types of Synchronization

1. **Synchronized Methods:** Entire methods are synchronized.
  - ▶ Instance methods: Lock on the instance.
  - ▶ Static methods: Lock on the class object.
2. **Synchronized Blocks:** Finer-grained control by synchronizing specific blocks of code.

# Synchronized Method

```
1  class Counter {  
2      private int count = 0;  
3      public synchronized void increment() {  
4          count++;  
5      }  
6  }
```

- ▶ Here, the `increment()` method is synchronized, ensuring mutual exclusion.
- ▶ Only one thread can execute `increment()` at a time.

# Synchronized Block

```
1  class Counter {  
2      private int count = 0;  
3      public void increment() {  
4          synchronized (this) {  
5              count++;  
6          }  
7      }  
8  }
```

- ▶ Here, only the critical section (`count++`) is synchronized.
- ▶ Non-critical operations in the same method can execute concurrently.

# Synchronized Static Method

```
1  class StaticCounter {  
2      private static int count = 0;  
3      public static synchronized void increment() {  
4          count++;  
5      }  
6  }
```

- ▶ The `increment()` method locks the class object, ensuring mutual exclusion across all threads accessing static data.

# Notes and Best Practices

- ▶ Synchronize only the code that truly needs to be thread-safe.
- ▶ Be mindful of **deadlocks**, where two threads are waiting for each other to release locks.
- ▶ Avoid synchronizing large blocks of code as it can degrade performance.

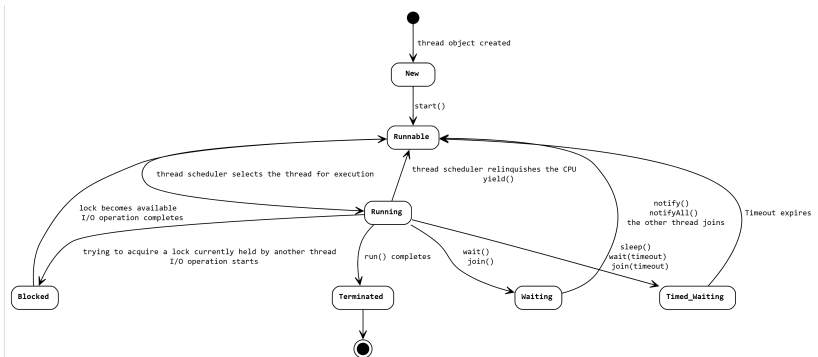
# Key Takeaways

- ▶ Use `synchronized` to protect shared resources in multithreaded environments.
- ▶ Understand the difference between synchronizing methods and blocks.
- ▶ Be cautious of potential performance issues and deadlocks.
- ▶ Always design for thread safety when shared mutable state is involved.

# Thread Lifecycle in Java

- ▶ A Java thread goes through multiple states during its lifecycle:
  - ▶ **NEW**: Thread is created but not started.
  - ▶ **RUNNABLE**: Ready to run but waiting for CPU time
  - ▶ **RUNNING**: Actively executing.
  - ▶ **BLOCKED**: Waiting for resources
  - ▶ **WAITING**: Waiting for a signal from another thread
  - ▶ **TIMED\_WAITING**: Waiting for timer to expire
  - ▶ **TERMINATED**: Finished execution.

# Thread Lifecycle





# Methods Controlling Thread Lifecycle

- ▶ `start()` - Starts the thread. Moves thread from NEW to RUNNABLE state.
- ▶ `run()` - Entry point of the thread. Defines the thread's behavior.
- ▶ `sleep(ms)` - Puts the thread to sleep. Moves thread to TIMED\_WAITING state.
- ▶ `join()` - Causes the current thread to wait for another thread to finish.
- ▶ `wait()`, `notify()`, `notifyAll()` - For inter-thread communication.
- ▶ `yield()` - Temporarily pauses the thread to allow others to execute.
- ▶ `interrupt()` - Interrupts a thread in certain states.

# start() and run()

start():

- ▶ Transitions thread from NEW to RUNNABLE.
- ▶ Calls run() internally.

run():

- ▶ Defines the thread's task.
- ▶ Invoked by the JVM when the thread starts.

## Example:

```
1  class MyThread extends Thread {  
2      public void run() {  
3          System.out.println("Thread is running");  
4      }  
5  }  
6  
7  public class Main {  
8      public static void main(String[] args) {  
9          Thread t = new MyThread();  
10         t.start(); // Calls run()  
11     }  
12 }
```

# sleep()

- ▶ Pauses the thread for a specified duration.
- ▶ Transitions thread to TIMED\_WAITING state.

## Example:

```
1  class SleepExample extends Thread {
2      public void run() {
3          try {
4              System.out.println("Thread sleeping for 2 seconds");
5              Thread.sleep(2000); // Sleep for 2 seconds
6              System.out.println("Thread awake!");
7          } catch (InterruptedException e) {
8              System.out.println("Thread interrupted");
9          }
10     }
11 }
12
13 public class Main {
14     public static void main(String[] args) {
15         new SleepExample().start();
16     }
17 }
```

# join()

- ▶ Causes the current thread to wait for another thread to finish.
- ▶ Useful for coordinating threads.

## Example:

```
1  class JoinExample extends Thread {
2      public void run() {
3          System.out.println("Thread running...");
4      }
5  }
6
7  public class Main {
8      public static void main(String[] args) {
9          Thread t = new JoinExample();
10         t.start();
11         try {
12             t.join(); // Wait for t to finish
13             System.out.println("Main thread continues");
14         } catch (InterruptedException e) {
15             e.printStackTrace();
16         }
17     }
18 }
```

# yield()

- ▶ Suggests that the thread scheduler give another thread an opportunity to execute.
- ▶ The scheduler may or may not honor this request.
- ▶ May move the thread back to RUNNABLE state.

## Example:

```
1  class YieldExample extends Thread {
2      public void run() {
3          for (int i = 0; i < 5; i++) {
4              System.out.println(Thread.currentThread().getName() + " - " + i);
5              Thread.yield(); // Hint scheduler to switch
6          }
7      }
8  }
9
10 public class Main {
11     public static void main(String[] args) {
12         new YieldExample().start();
13         new YieldExample().start();
14     }
15 }
```

# interrupt()

- ▶ Sends an interrupt signal to a thread.
- ▶ Does not forcibly stop the thread but sets an interrupt flag.
- ▶ Commonly used to stop threads in WAITING, TIMED\_WAITING, or BLOCKED state. Throws an InterruptedException.

## Example:

```
1  class InterruptExample extends Thread {
2      public void run() {
3          try {
4              Thread.sleep(5000); // Sleep for 5 seconds
5          } catch (InterruptedException e) {
6              System.out.println("Thread interrupted");
7          }
8      }
9  }
10
11 public class Main {
12     public static void main(String[] args) {
13         Thread t = new InterruptExample();
14         t.start();
15         t.interrupt(); // Interrupt the thread
16     }
17 }
```

## wait(), notify(), notifyAll()

- ▶ Used for inter-thread communication.
- ▶ wait() pauses a thread until another thread invokes notify() or notifyAll().
- ▶ Must be used in a synchronized block or method.

### Example:

```
1  class Shared {
2      synchronized void produce() throws InterruptedException {
3          System.out.println("Producing...");
4          wait(); // Wait for notification
5          System.out.println("Resumed production");
6      }
7      synchronized void consume() {
8          System.out.println("Consuming...");
9          notify(); // Notify waiting thread
10     }
11 }
12 public class Main {
13     public static void main(String[] args) {
14         Shared shared = new Shared();
15         new Thread(() -> {
16             try {
17                 shared.produce();
18             } catch (InterruptedException e) { }
19         }).start();
20         new Thread(shared::consume).start();
21     }
22 }
```

# wait(), notify(), and notifyAll() in Java

- ▶ These methods are part of the 'Object' class.
- ▶ Used for **inter-thread communication** within a synchronized context.
- ▶ Allows one thread to pause execution and wait until another thread signals it to continue.



# Why wait(), notify(), and notifyAll() are in the Object Class

- ▶ **Object-Level Locking:**

- ▶ Every Java object has an intrinsic lock (or monitor).
- ▶ wait(), notify(), and notifyAll() operate on this lock.

- ▶ **Synchronization is Object-Oriented:**

- ▶ Multiple threads may coordinate using a shared object.
- ▶ The shared object is the entity that ensures proper thread communication.

- ▶ **Thread Communication:**

- ▶ wait(), notify(), and notifyAll() are designed for inter-thread communication via shared resources.
- ▶ These methods naturally belong to the object being shared, not the thread itself.

# wait()

- ▶ **Definition:** Makes the current thread wait until another thread invokes 'notify()' or 'notifyAll()' on the same object.
- ▶ Must be called inside a `synchronized` block or method.
- ▶ Releases the lock on the object temporarily, allowing other threads to acquire it.
- ▶ Syntax:

```
1  synchronized (object) {  
2      while (condition) {  
3          object.wait();  
4      }  
5  }
```

# notify()

- ▶ **Definition:** Wakes up a single thread that is waiting on the object's monitor.
  - ▶ The thread to wake up is chosen by the JVM scheduler (not guaranteed to be FIFO).
- ▶ The thread awakened must re-acquire the lock before resuming execution.
- ▶ Syntax:

```
1 synchronized (object) {  
2     object.notify();  
3 }
```

- ▶ Useful when only one waiting thread needs to be notified.

# notifyAll()

- ▶ **Definition:** Wakes up all threads that are waiting on the object's monitor.
- ▶ Threads compete to acquire the lock and proceed based on thread scheduling.

- ▶ **Syntax:**

```
1 synchronized (object) {  
2     object.notifyAll();  
3 }
```

- ▶ Useful in scenarios where all waiting threads must proceed.

## Rules for Using `wait()`, `notify()`, `notifyAll()`

- ▶ Must be called from within a `synchronized` block or method.
- ▶ Always ensure a proper condition check (e.g., `while` loop) when using `wait()`.
- ▶ Avoid relying on thread priorities; behavior depends on the JVM scheduler.
- ▶ Use `notifyAll()` in complex scenarios to avoid deadlocks.
- ▶ Avoid "spurious wakeups" by rechecking conditions after `wait()`.

## Example: Producer-Consumer with wait() and notify()

```
1  // Shared object
2  class Queue {
3      private List<Integer> buffer = new ArrayList<>();
4      private final int SIZE = 5;
5
6      public synchronized void produce(int value) throws InterruptedException {
7          while (buffer.size() == SIZE) wait(); // Try using if
8          buffer.add(value);
9          notifyAll(); // Notify consumers
10     }
11
12     public synchronized int consume() throws InterruptedException {
13         while (buffer.isEmpty()) wait();
14         int value = buffer.remove(0);
15         notifyAll(); // Notify producers
16         return value;
17     }
18 }
```

# When to Use `notify()` vs `notifyAll()`

- ▶ Use `notify()` when:
  - ▶ Only one waiting thread should proceed.
  - ▶ There is a single condition to check.
- ▶ Use `notifyAll()` when:
  - ▶ Multiple threads may be waiting for different conditions.
  - ▶ You want to avoid potential deadlocks.
- ▶ Example:
  - ▶ A producer-consumer system with multiple producers and consumers benefits from `notifyAll()`.

## Exercise: Sum of Array

- ▶ Write a multithreaded program to calculate the sum of an array using:
  - ▶ Multiple threads
  - ▶ Proper synchronization